

Practical 03 — Choosing a Model Honestly

SCSE3040 Machine Learning Operations · Bennett University · Roll number 24f2000604

Practical 03 --- Choosing a Model Honestly

Overfitting, the train/test gap, and cross-validation

SCSE3040 Machine Learning Operations · Bennett University · Session 2026-27

Lectures this follows	L04-L05
Course Outcome	CO2
Lab duration	120 minutes
Memory needed	about 400 MB (fine on a 4 GB or 8 GB machine)
Extra software	nothing beyond the course venv
Marks	10

Aim

1. Train three different models on the same data and compare them fairly.
 2. See overfitting happen, and measure it as a gap between two scores.
 3. Use cross-validation instead of trusting one lucky split.
 4. Choose a model for a reason you can defend, not because it scored best.
-

What you need to know first

In P02 you trained one model and it did well. A reasonable next thought is "let me try a more powerful model and do even better". Today you find out why that thought is dangerous.

A **decision tree** is a model that asks yes/no questions: *is the distance more than 6 km? is it raining?* Ask enough questions and you can separate every single training order from every other one. At that point the model has not learned a pattern --- it has memorised the training data. It will look perfect in training and fall apart on anything new.

That failure has a name: **overfitting**. And it has a measurement. Score the model twice, once on the data it trained on and once on data it has never seen. The distance between those two numbers is the **train/test gap**. A small gap means the model learned something general. A big gap means it memorised.

There is one more problem to fix. In P02 your test set was one particular 120 orders, chosen by one particular seed. What if those 120 happened to be easy? **Cross-validation** removes that luck: it cuts the data into 5 equal parts, trains 5 times, and each time tests on a different part. You get 5 scores instead of 1, and you report the average. It costs 5 times the computing time, which on this dataset means well under a second.

The habit to build today: *the best score is not the best model*. A simple model that scores almost as well is usually the one to ship, because you can explain it, it trains faster, and it breaks less.

Before you start

- Practical P02 is finished. You can split data, train a model and read an MAE.
- You remember that the test set is the only honest score.

How to run a cell

Click on a grey code cell, then press **Shift + Enter**. The cell runs and the cursor moves to the next one. A number appears in the [] on the left when the cell has finished.

Run the cells in order, from the top. A later cell almost always uses something an earlier cell created. If you jump ahead you will see a `NameError`, which just means "you have not made that thing yet".

If everything goes wrong, use the menu: **Kernel -> Restart Kernel and Clear All Outputs**, then start again from the first cell. Nothing is damaged by doing this.

In [1]:

```
# Step 0 --- check the workbench before we start.
# This cell only looks; it changes nothing. Run it and read the last line.

import sys
from pathlib import Path

print("Python  :", sys.version.split()[0])
print("Folder  :", Path.cwd().name)

_missing = []
for _name in ['numpy', 'pandas', 'sklearn', 'matplotlib']:
    try:
        __import__(_name)
    except ImportError:
        _missing.append(_name)

for _name in ['numpy', 'pandas', 'sklearn', 'matplotlib']:
    _mark = "missing" if _name in _missing else "ok"
    print(f"  {_name:<14} {_mark}")

if _missing:
    print()
    print("STOP. Some libraries are missing:", ", ".join(_missing))
    print("Ask your instructor to run the setup in labs/SETUP.md.")
else:
    print()
    print("All good. You can carry on to Step 1.")
```

```
Python  : 3.13.12
Folder   : P03-model-choice
```

```
numpy      ok
pandas     ok
sklearn    ok
matplotlib ok
```

All good. You can carry on to Step 1.

Step 0b --- the dataset

Every practical in this course uses the same 600 food deliveries. The next cell makes sure the file is there.

In [2]:

```
# The delivery-time dataset every practical in this course uses.
# If the file is missing we build it again from the same seed, so every
# student in the room gets byte-for-byte the same 600 rows.

import csv
from pathlib import Path

import numpy as np

SEED = 42
N_ROWS = 600
DATA = Path("..") / "data" / "delivery_times.csv"

def make_delivery_csv(path=DATA):
    """Write the 600-row delivery dataset. Same formula as the lectures."""
    rng = np.random.default_rng(SEED)
    distance_km = np.round(rng.uniform(0.5, 12.0, N_ROWS), 2)
    prep_time_min = np.round(rng.uniform(5, 30, N_ROWS), 0)
    traffic_level = rng.integers(1, 4, N_ROWS)
```

```

rain = rng.binomial(1, 0.25, N_ROWS)
delivery_min = np.round(
    6.0
    + 3.1 * distance_km
    + 0.65 * prep_time_min
    + 4.2 * traffic_level
    + 5.5 * rain
    + rng.normal(0, 2.5, N_ROWS),
    1,
)
path.parent.mkdir(parents=True, exist_ok=True)
with path.open("w", newline="", encoding="utf-8") as fh:
    w = csv.writer(fh)
    w.writerow(["distance_km", "prep_time_min", "traffic_level",
                "rain", "delivery_min"])
    for i in range(N_ROWS):
        w.writerow([distance_km[i], int(prep_time_min[i]),
                    int(traffic_level[i]), int(rain[i]), delivery_min[i]])
return path

if not DATA.exists():
    make_delivery_csv()
    print("dataset rebuilt ->", DATA)
else:
    print("dataset found ->", DATA)

```

dataset found -> ../data/delivery_times.csv

Walkthrough

Read each step, then run its cell.

Step 1 --- Set up the data once, for all three models

Same data, same split, same seed as P02. Every model today gets exactly the same deal --- otherwise the comparison is worthless.

In [3]:

```

import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split

orders = pd.read_csv(DATA)

FEATURES = ["distance_km", "prep_time_min", "traffic_level", "rain"]
X = orders[FEATURES]
y = orders["delivery_min"]

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42
)

print("train:", len(X_train), " test:", len(X_test))

```

train: 480 test: 120

Step 2 --- A helper that scores a model both ways

We will do the same two measurements many times today, so write them once. This function returns the training MAE, the test MAE, and the gap between them.

In [4]:

```

from sklearn.metrics import mean_absolute_error

def score_both_ways(model, name):
    """Train the model, then score it on train and on test."""
    model.fit(X_train, y_train)
    train_mae = mean_absolute_error(y_train, model.predict(X_train))

```

```

test_mae = mean_absolute_error(y_test, model.predict(X_test))
gap = test_mae - train_mae
print(f"{name:<26} train {train_mae:5.2f}      "
      f"test {test_mae:5.2f}      gap {gap:5.2f}")
return {"name": name, "train": train_mae,
        "test": test_mae, "gap": gap}

print("helper ready")

```

helper ready

Step 3 --- Model 1: linear regression, the honest baseline

The model from P02. Watch the gap: it should be tiny, because a straight line is too simple to memorise anything.

In [5]:

```

from sklearn.linear_model import LinearRegression

linear = score_both_ways(LinearRegression(), "LinearRegression")

```

```
LinearRegression      train  2.04    test  1.92    gap -0.11
```

Train and test scores almost equal. This model is not overfitting. It may not be the most accurate thing possible, but it is telling you the truth.

Step 4 --- Model 2: a decision tree with no limits

Now the dangerous one. A decision tree with no maximum depth is allowed to keep asking questions until every training order sits in its own little box.

Look at the training score first. Then look at the test score.

In [6]:

```

from sklearn.tree import DecisionTreeRegressor

wild_tree = score_both_ways(
    DecisionTreeRegressor(random_state=42), "DecisionTree (no limit)"
)

```

```
DecisionTree (no limit)  train  0.00    test  3.43    gap  3.43
```

A training MAE near **zero** and a much worse test MAE. The model got every training order perfectly right and learned nothing useful.

That gap is overfitting. You have now seen it happen, not just been told about it.

Step 5 --- Model 3: the same tree, but shallow

Now stop the tree early. `max_depth=4` means it may only ask four questions before it has to answer. It can no longer memorise.

In [7]:

```

small_tree = score_both_ways(
    DecisionTreeRegressor(max_depth=4, random_state=42),
    "DecisionTree (depth 4)",
)

```

```
DecisionTree (depth 4)  train  3.66    test  4.23    gap  0.57
```

Worse on training, better on test, and a far smaller gap. Limiting a model's freedom made it *more* useful. That idea is the whole of model selection.

Step 6 --- Model 4: a random forest

A **random forest** grows many different trees and averages their answers. Individual trees still overfit, but their mistakes point in different directions and largely cancel out.

We use 50 trees. More would be slower and no better here.

In [8]:

```
from sklearn.ensemble import RandomForestRegressor

forest = score_both_ways(
    RandomForestRegressor(n_estimators=50, random_state=42),
    "RandomForest (50 trees)",
)

RandomForest (50 trees)    train  1.00    test  2.39    gap  1.39
```

Step 7 --- Put the four results side by side

Never compare models from memory. Put them in a table.

In [9]:

```
results = pd.DataFrame([linear, wild_tree, small_tree, forest])
results = results.round(2).sort_values("test")

print(results.to_string(index=False))
print()
best = results.iloc[0]
print(f"Lowest test MAE: {best['name']} at {best['test']:.2f} minutes")
```

	name	train	test	gap
	LinearRegression	2.04	1.92	-0.11
	RandomForest (50 trees)	1.00	2.39	1.39
	DecisionTree (no limit)	0.00	3.43	3.43
	DecisionTree (depth 4)	3.66	4.23	0.57

Lowest test MAE: LinearRegression at 1.92 minutes

Read the `gap` column as carefully as the `test` column. A model with a good test score and a huge gap got lucky, and luck does not survive contact with real users.

Step 8 --- Draw the gap, because a table makes you read

The table is correct but slow --- you must read eight numbers and subtract in your head. A chart does the subtraction for your eyes.


matplotlib is Python's standard charting library. It has been sitting in your course environment since P01; today it earns its place. The first line, `%matplotlib inline`, tells Jupyter to draw the picture right below the cell.

In [10]:

```
%matplotlib inline
import matplotlib.pyplot as plt

four = [linear, wild_tree, small_tree, forest]
xs = np.arange(len(four))

fig, ax = plt.subplots(figsize=(9, 3.8))
ax.bar(xs - 0.2, [m["train"] for m in four], 0.38,
       label="train MAE", color="#6B7785")
ax.bar(xs + 0.2, [m["test"] for m in four], 0.38,
       label="test MAE", color="#9E3B33")
ax.set_xticks(xs, [m["name"] for m in four], fontsize=9)
ax.set_ylabel("MAE (minutes)")
ax.set_title("The distance between a pair of bars is the gap")
ax.legend()
plt.tight_layout()
plt.show()
```



Look at the pair for the no-limit tree. Its grey train bar has **vanished** --- an MAE of 0.00 draws as no bar at all --- while its red test bar is among the tallest. Bars far apart mean the model memorised; bars close together mean it learned.

The linear model's two bars are the closest pair on the chart. That is what honesty looks like.

Step 9 --- Stop trusting one split: cross-validation

Every number above came from one 120-order test set. Change the seed and they all move a little.

5-fold cross-validation cuts the data into 5 parts and trains 5 times, testing on a different part each time. We report the average and the spread.

scikit-learn reports errors as *negative* numbers, because its convention is that bigger is better. We flip the sign back.

In [11]:

```
from sklearn.model_selection import cross_val_score

def cross_validate(model, name):
    """Average MAE over 5 different train/test splits."""
    scores = -cross_val_score(
        model, X, y, cv=5, scoring="neg_mean_absolute_error"
    )
    print(f"{name:<26} MAE {scores.mean():5.2f} "
          f" (+/- {scores.std():4.2f})")
    return scores.mean()

cv_linear = cross_validate(LinearRegression(), "LinearRegression")
cv_tree = cross_validate(
    DecisionTreeRegressor(max_depth=4, random_state=42),
    "DecisionTree (depth 4)"
)
cv_forest = cross_validate(
    RandomForestRegressor(n_estimators=50, random_state=42),
    "RandomForest (50 trees)"
)
```

```
LinearRegression      MAE  2.03 (+/- 0.16)
DecisionTree (depth 4) MAE  4.57 (+/- 0.18)

RandomForest (50 trees) MAE  2.69 (+/- 0.40)
```

The +/- is the spread across the 5 folds. If two models are closer together than their spread, you cannot honestly say one is better.

Step 10 --- Tune one setting, properly

How deep should the tree be? Do not guess, and do not try depths on the test set --- that would leak the answers.

Try each depth with cross-validation, which never touches the test set at all.

In [12]:

```
depths = [2, 3, 4, 6, 8, None]
scores_by_depth = {}

for depth in depths:
    tree = DecisionTreeRegressor(max_depth=depth, random_state=42)
    mae = -cross_val_score(
        tree, X, y, cv=5, scoring="neg_mean_absolute_error"
    ).mean()
    scores_by_depth[depth] = round(mae, 3)
    print(f"max_depth={str(depth):<5} MAE {mae:5.2f}")

best_depth = min(scores_by_depth, key=scores_by_depth.get)
print()
print("best depth:", best_depth)
```

```
max_depth=2      MAE  5.86
max_depth=3      MAE  4.85
max_depth=4      MAE  4.57
max_depth=6      MAE  3.64
max_depth=8      MAE  3.39
max_depth=None   MAE  3.48
```

```
best depth: 8
```

Read the **shape**, not just the winner.

Depth 2 is the worst of all: too simple to capture anything, which is called **underfitting**. The scores improve as the tree is allowed more room, reach their best in the middle, and then get slightly worse again when the limit is removed altogether.

Too rigid and too free are both bad. The good answer is in between, and the only way to find it is to measure --- which is what you just did, without ever touching the test set.

Step 11 --- Draw the sweep, because ten numbers hide a shape


The loop printed six scores. Their *shape* is the real lesson, and a shape is easier to see than to read. Draw MAE against depth.

The red circle marks the depth cross-validation liked best.

In [13]:

```
xs = np.arange(len(depths))
maes = [scores_by_depth[d] for d in depths]
best_i = int(np.argmin(maes))

fig, ax = plt.subplots(figsize=(8, 3.6))
ax.plot(xs, maes, marker="o", color="#2C5F79", linewidth=2)
ax.plot(xs[best_i], maes[best_i], marker="o", markersize=13,
        markerfacecolor="none", markeredgecolor="#9E3B33",
        markeredgewidth=2)
ax.annotate(f"best: depth {depths[best_i]}",
            (xs[best_i], maes[best_i]),
            xytext=(0, -22), textcoords="offset points",
            ha="center", color="#9E3B33")
ax.set_xticks(xs, [str(d) for d in depths])
ax.set_ylim(bottom=min(maes) - 0.8)
ax.set_xlabel("max_depth")
ax.set_ylabel("cross-validated MAE (minutes)")
ax.set_title("Too rigid on the left, too free on the right")
plt.tight_layout()
plt.show()
```



Now the shape is unmissable. The left end towers over everything: depth 2 is the worst score on the whole chart --- **underfitting**, a model too rigid to learn anything. The curve falls as the tree earns more freedom, bottoms out in the middle, then ticks back up when the last limit is removed.

Both ends of the curve are bad. The answer lives at the bottom, and you found it without ever touching the test set.

Step 12 --- Choose, and be able to say why

Now make the decision the way a professional would, out loud.

In [14]:

```
print("Cross-validated MAE, best first:")
for name, mae in sorted(
    {"LinearRegression": cv_linear,
     "DecisionTree(4)": cv_tree,
     "RandomForest(50)": cv_forest}.items(),
    key=lambda kv: kv[1],
):
    print(f"  {name:<20} {mae:5.2f}")

print()
print("Our choice for the delivery service: LinearRegression.")
print("Reasons:")
print("  1. It has the lowest cross-validated error here.")
print("  2. It trains in milliseconds and needs almost no memory.")
print("  3. We can read its coefficients and explain a prediction")
print("     to a restaurant owner. Nobody can read 50 trees.")
print()
print("Be careful with reason 1. This dataset really was built")
print("from a straight line, so a straight line wins. On messier")
print("real data the forest usually wins instead -- and then you")
print("must decide what its extra accuracy is worth, against")
print("reasons 2 and 3, which do not change.")
```

```
Cross-validated MAE, best first:
LinearRegression      2.03
RandomForest(50)     2.69
DecisionTree(4)       4.57
```

```
Our choice for the delivery service: LinearRegression.
Reasons:
```

1. It has the lowest cross-validated error here.
2. It trains in milliseconds and needs almost no memory.
3. We can read its coefficients and explain a prediction to a restaurant owner. Nobody can read 50 trees.

```
Be careful with reason 1. This dataset really was built
```

from a straight line, so a straight line wins. On messier real data the forest usually wins instead -- and then you must decide what its extra accuracy is worth, against reasons 2 and 3, which do not change.

You will not always choose the simple model. But you must always be able to say why you chose what you chose. "It scored best" is not a reason a senior engineer will accept on its own --- and, as the last paragraph says, it will not always be true either.

Your turn

The walkthrough above is finished. Now you write some code.

There are **3 tasks**. Each one is small. Each one has a hint. Do them in order.

Where you see `# TODO`, replace that line with your own code. Do not delete the variable name on the left of the `=` sign --- the self-check at the bottom looks for exactly that name.

When you have tried all three, run the **self-check** cell at the end. It prints a table telling you which tasks are correct. You can run it as many times as you like.

Task T1 --- Measure overfitting yourself

Train a `DecisionTreeRegressor` with `max_depth=12` and `random_state=42`, and work out its train/test gap by hand.

Store the three numbers in `T1_train_mae`, `T1_test_mae` and `T1_gap`.

`T1_gap` must be the test MAE **minus** the train MAE.

Hint. You may not use `score_both_ways` for this one -- write the three lines out, so you can see what the helper was doing. Fit the tree, then call `mean_absolute_error` twice: once with `y_train` and `X_train`, once with `y_test` and `X_test`.

In [15]:

```
deep_tree = DecisionTreeRegressor(max_depth=12, random_state=42)

# Written out by hand rather than through score_both_ways, so the three
# steps the helper hides are visible.
deep_tree.fit(X_train, y_train)
T1_train_mae = mean_absolute_error(y_train, deep_tree.predict(X_train))
T1_test_mae = mean_absolute_error(y_test, deep_tree.predict(X_test))
T1_gap = T1_test_mae - T1_train_mae

print(f"train {T1_train_mae}  test {T1_test_mae}  gap {T1_gap}")

train 0.04695601851851851  test 3.397986111111111  gap 3.3510300925925924
```

Task T2 --- Tune the forest, without touching the test set

A random forest has a setting `max_depth` too. Use **5-fold cross-validation** to find which of `[2, 4, 6, None]` is best for a `RandomForestRegressor` (`n_estimators=50`, `random_state=42`).

Build a dictionary `T2_scores` mapping each depth to its cross-validated MAE, then put the winning depth in `T2_best_depth`.

You must not use `X_test` or `y_test` anywhere in this task.

Hint. Step 10 does exactly this for a single tree. Copy that loop and change two things: the model, and the list of depths. The winner is `min(T2_scores, key=T2_scores.get)`.

In [16]:

```
T2_depths = [2, 4, 6, None]

# Cross-validation only ever sees X and y -- the test set stays untouched.
T2_scores = {}
for depth in T2_depths:
```

```

forest = RandomForestRegressor(
    n_estimators=50, max_depth=depth, random_state=42
)
mae = -cross_val_score(
    forest, X, y, cv=5, scoring="neg_mean_absolute_error"
).mean()
T2_scores[depth] = round(mae, 3)
print(f"max_depth={str(depth):<5} MAE {mae:5.2f}")

T2_best_depth = min(T2_scores, key=T2_scores.get)

print(T2_scores)
print("best:", T2_best_depth)

max_depth=2      MAE   5.27

max_depth=4      MAE   3.64

max_depth=6      MAE   2.84

max_depth=None   MAE   2.69
{2: np.float64(5.271), 4: np.float64(3.643), 6: np.float64(2.844), None: np.float64(2.688)}
best: None

```

Task T3 --- A comparison you can reuse

Write a function `compare(models)` that takes a **dictionary** of `{name: model}` and returns a dictionary of `{name: test MAE}`.

It must fit each model on `X_train/y_train` and score it on `X_test/y_test`.

Then call it on the three models given in the stub and store the answer in `T3_table`.

Hint. Loop over `models.items()`, call `.fit(...)`, then `mean_absolute_error(y_test, m.predict(X_test))`, and collect the results into a new dictionary that you `return` at the end.

In [17]:

```

def compare(models):
    """Fit every model on the training data, score each on the test set."""
    results = {}
    for name, model in models.items():
        model.fit(X_train, y_train)
        results[name] = mean_absolute_error(y_test, model.predict(X_test))
    return results

T3_table = compare({
    "linear": LinearRegression(),
    "tree4": DecisionTreeRegressor(max_depth=4, random_state=42),
    "forest": RandomForestRegressor(n_estimators=50, random_state=42),
})

print(T3_table)

{'linear': 1.924653079833835, 'tree4': 4.228332168713071, 'forest': 2.3918166666666663}

```

Self-check

Run the cell below to mark your work.

In [18]:

```

# -----
# SELF-CHECK --- run this when you have attempted the tasks above.
# It never breaks your notebook. A task you have not done yet simply
# shows FAIL.
# -----

```

```

_results = []

def _check(label, fn):
    """Evaluate one graded condition without ever raising."""
    try:
        ok = bool(fn())
    except Exception:
        ok = False
    _results.append((label, ok))

_check('T1 | the depth-12 tree is trained', lambda: hasattr(deep_tree, 'tree_'))
_check('T1 | T1_train_mae and T1_test_mae are that tree's two scores', lambda: abs(float(T1_train
_check('T1 | T1_gap is test minus train, and it is positive', lambda: abs(float(T1_gap) - (float(
_check('T2 | T2_scores has one score per depth', lambda: set(T2_scores) == {2, 4, 6, None})
_check('T2 | every score is a plausible MAE in minutes', lambda: all(0 < float(v) < 20 for v in T
_check('T2 | T2_best_depth is the depth with the lowest score', lambda: T2_best_depth == min(T2_s
_check('T3 | compare() returns one score per model', lambda: set(T3_table) == {'linear', 'tree4',
_check('T3 | the linear score matches the walkthrough', lambda: abs(float(T3_table['linear']) - 1

print("=====")
print("SELF-CHECK   Practical 03 --- Choosing a Model Honestly")
print("=====")
for _label, _ok in _results:
    print(f"   [{'PASS' if _ok else 'FAIL'}]   {_label}")
print("-----")
_passed = sum(1 for _, _ok in _results if _ok)
print(f"   {_passed} of {len(_results)} checks passed")
print("=====")
if _passed == len(_results):
    print("Well done. Save the notebook and submit it.")
else:
    print("Read the FAIL lines above, fix those tasks, run this cell again.")

=====
SELF-CHECK   Practical 03 --- Choosing a Model Honestly
=====
[PASS] T1 | the depth-12 tree is trained
[PASS] T1 | T1_train_mae and T1_test_mae are that tree's two scores
[PASS] T1 | T1_gap is test minus train, and it is positive
[PASS] T2 | T2_scores has one score per depth
[PASS] T2 | every score is a plausible MAE in minutes
[PASS] T2 | T2_best_depth is the depth with the lowest score
[PASS] T3 | compare() returns one score per model
[PASS] T3 | the linear score matches the walkthrough
-----
8 of 8 checks passed
=====
Well done. Save the notebook and submit it.

```

Model choice --- three sentences

I would deploy plain `LinearRegression()` fitted on the four features `distance_km`, `prep_time_min`, `traffic_level` and `rain`, with no tuned hyperparameters, in preference to `RandomForestRegressor(n_estimators=50)`.

I choose it because 5-fold cross-validation scores it at **MAE 2.03 +/- 0.16 minutes** against **2.69 +/- 0.40** for the 50-tree forest and **4.57 +/- 0.18** for the depth-4 tree, and its train/test gap of **-0.11 min** (versus 1.39 for the forest and 3.43 for the unlimited tree) shows the score is honest rather than memorised; operationally it retrains in milliseconds, ships as four coefficients and an intercept instead of 50 stored trees, and lets me tell a restaurant partner exactly which term pushed a quote to 32 minutes.

The trade-off I accept is capacity: a straight line cannot represent the interactions a forest can --- rain making heavy traffic disproportionately worse, for example --- so on messier real orders than this generated dataset the forest's extra flexibility would likely win, and I would have to re-run this comparison and switch if monitoring showed the linear residuals bending.

What to submit

1. This notebook, with every cell run and its output visible.
2. In a markdown cell at the end, three sentences: which model you would put in front of real customers, why, and what you gave up by choosing it.

Name your file `PXX_<your-roll-number>.ipynb` before you upload it.

How this practical is marked

What is marked	Marks
Walkthrough run end to end, outputs visible	3
Task T1 --- overfitting measured as a gap	2
Task T2 --- best setting found by cross-validation	3
Task T3 --- a reusable comparison function	2
Total	10

Read more

- scikit-learn --- Underfitting vs overfitting --- https://scikit-learn.org/stable/auto_examples/model_selection/plot_underfitting_overfitting.html
- scikit-learn --- Cross-validation --- https://scikit-learn.org/stable/modules/cross_validation.html
- scikit-learn --- Decision trees --- <https://scikit-learn.org/stable/modules/tree.html>
- scikit-learn --- Ensemble methods --- <https://scikit-learn.org/stable/modules/ensemble.html>
- matplotlib --- Quick start guide --- https://matplotlib.org/stable/users/explain/quick_start.html